

Lesson 05 Demo 05

Building an Automated Log Processing Pipeline with Kafka and ELK Stack

Objective: To integrate Kafka with ELK Stack to buffer and process log surges, protecting Logstash and Elasticsearch during high-volume events for improved reliability and scalable log management

Tools required: ELK Stack and Kafka

Prerequisites: None

Steps to be followed:

1. Set up the system and install Java Runtime Environment (JRE)
2. Set up the Apache Kafka message broker
3. Create a topic for the Apache logs

Step 1: Set up the system and install Java Runtime Environment (JRE)

- 1.1 Run the following commands in the Ubuntu lab to log in as the root user, gain admin access, and update Ubuntu packages:

sudo su -
apt update



```
Applications root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
labuser@ip-172-31-34-121:~$ sudo su -
root@ip-172-31-34-121:~#
root@ip-172-31-34-121:~#
root@ip-172-31-34-121:~# apt update
```

```
Applications root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
labuser@ip-172-31-34-121:~$ sudo su -
root@ip-172-31-34-121:~#
root@ip-172-31-34-121:~#
root@ip-172-31-34-121:~# apt update
Hit:1 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy InRelease
Hit:2 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy-updates InRelease
Hit:3 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy-backports InRelease
Hit:4 https://artifacts.elastic.co/packages/7.x/apt stable InRelease
Hit:5 https://artifacts.elastic.co/packages/8.x/apt stable InRelease
Hit:6 https://download.docker.com/linux/ubuntu jammy InRelease
Ign:7 https://pkg.jenkins.io/debian-stable binary/ InRelease
Hit:8 https://apt.grafana.com stable InRelease
Hit:9 https://pkg.jenkins.io/debian-stable binary/ Release
Get:11 http://ports.ubuntu.com/ubuntu-ports jammy-security InRelease [129 kB]
Hit:10 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.30/deb InRelease
Hit:12 https://ppa.launchpadcontent.net/ansible/ansible/ubuntu jammy InRelease
Hit:13 https://ppa.launchpadcontent.net/pipewire-debian/pipewire-upstream/ubuntu jammy InRelease
Fetched 129 kB in 2s (67.7 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
152 packages can be upgraded. Run 'apt list --upgradable' to see them.
W: https://artifacts.elastic.co/packages/8.x/apt/dists/stable/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
W: https://download.docker.com/linux/ubuntu/dists/jammy/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
root@ip-172-31-34-121:~#
```

1.2 Execute the following command to install JRE, which Filebeat will use:

apt -y install default-jre

```
root@ip-172-31-43-84:~# apt -y install default-jre
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  default-jre-headless
The following NEW packages will be installed:
  default-jre default-jre-headless
0 upgraded, 2 newly installed, 0 to remove and 267 not upgraded.
Need to get 3938 B of archives.
After this operation, 26.6 kB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy/main amd64 default-jre-headless amd64 2:1.11-72build2 [3042 B]
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy/main amd64 default-jre amd64 2:1.11-72build2 [896 B]
Fetched 3938 B in 0s (222 kB/s)
Selecting previously unselected package default-jre-headless.
(Reading database ... 50%
```

After JRE is installed, proceed with the installation of the Elasticsearch component of ELK Stack. To do this, add the signing key and repositories to the system first.

- 1.3 Execute the following command to add the Elasticsearch signing key:
- ```
wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo gpg --dearmor -o /usr/share/keyrings/elasticsearch-keyring.gpg
```



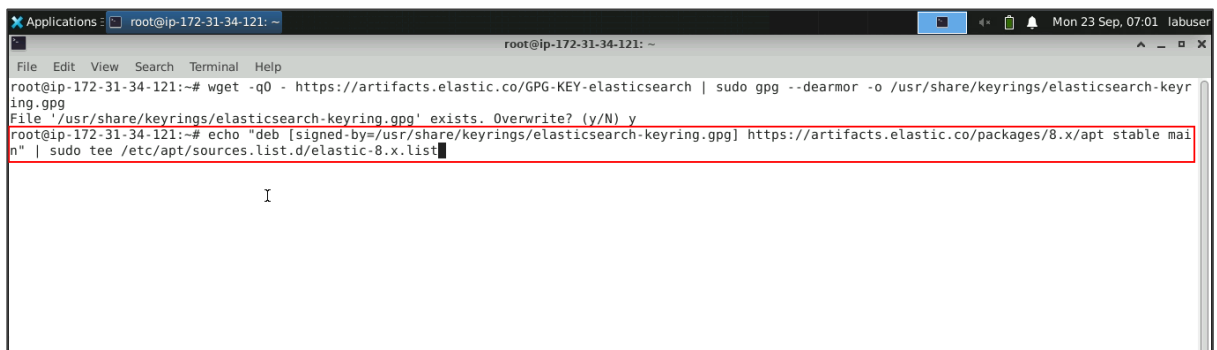
```
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo gpg --dearmor -o /usr/share/keyrings/elasticsearch-keyring.gpg
```



```
root@ip-172-31-34-121:~# wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo gpg --dearmor -o /usr/share/keyrings/elasticsearch-keyring.gpg
File '/usr/share/keyrings/elasticsearch-keyring.gpg' exists. Overwrite? (y/N) y
```

**Note:** If the command displays an error indicating that the file already exists, type **y** and press the **enter** key to proceed with overwriting the existing file.

- 1.4 Run the following command to add the repository in `/etc/apt/sources.list.d/elastic-8.x.list`:
- ```
echo "deb [signed-by=/usr/share/keyrings/elasticsearch-keyring.gpg]  
https://artifacts.elastic.co/packages/8.x/apt stable main" | sudo tee  
/etc/apt/sources.list.d/elastic-8.x.list
```



```
root@ip-172-31-34-121:~# wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo gpg --dearmor -o /usr/share/keyrings/elasticsearch-keyring.gpg  
File '/usr/share/keyrings/elasticsearch-keyring.gpg' exists. Overwrite? (y/N) y  
root@ip-172-31-34-121:~# echo "deb [signed-by=/usr/share/keyrings/elasticsearch-keyring.gpg] https://artifacts.elastic.co/packages/8.x/apt stable main" | sudo tee /etc/apt/sources.list.d/elastic-8.x.list
```

1.5 Run following command to update the system packages:

apt update

```
root@ip-172-31-36-251:~#
root@ip-172-31-36-251:~# wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo gpg --dearmor -o /usr/share/keyrings/elasticsearch-keyring.gpg
root@ip-172-31-36-251:~#
root@ip-172-31-36-251:~# echo "deb [signed-by=/usr/share/keyrings/elasticsearch-keyring.gpg] https://artifacts.elastic.co/packages/8.x/apt stable main" | sudo tee /etc/
apt/sources.list.d/elastic-8.x.list
deb [signed-by=/usr/share/keyrings/elasticsearch-keyring.gpg] https://artifacts.elastic.co/packages/8.x/apt stable main
root@ip-172-31-36-251:~#
root@ip-172-31-36-251:~# apt update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-backports InRelease
Get:4 https://artifacts.elastic.co/packages/8.x/apt stable InRelease [10.4 kB]
Hit:5 https://download.docker.com/linux/ubuntu jammy InRelease
Ign:6 https://pkg.jenkins.io/debian-stable binary/ InRelease
Hit:7 https://pkg.jenkins.io/debian-stable binary/ Release
Hit:9 http://security.ubuntu.com/ubuntu jammy-security InRelease
Hit:8 https://prod-cdn.packages.k8s.io/repositories/istio/kubernetes:/core:/stable:/v1.28/deb InRelease
Get:10 https://artifacts.elastic.co/packages/8.x/apt stable/main amd64 Packages [109 kB]
Fetched 109 kB in 1s (155 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
256 packages can be upgraded. Run 'apt list --upgradable' to see them.
W: Ignoring file 'elastic-8.x.listecho' in directory '/etc/apt/sources.list.d/' as it has an invalid filename extension
root@ip-172-31-36-251:~#
```

1.6 Run the following command to configure a Logstash pipeline that pulls logs from a Kafka topic, processes them, and ships them to Elasticsearch for indexing:

vim /etc/logstash/conf.d/apache.conf

```
Applications root@ip-172-31-34-121: ~
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo gpg --dearmor -o /usr/share/keyrings/elasticsearch-keyr
ing.gpg
File '/usr/share/keyrings/elasticsearch-keyring.gpg' exists. Overwrite? (y/N) y
root@ip-172-31-34-121:~# echo "deb [signed-by=/usr/share/keyrings/elasticsearch-keyring.gpg] https://artifacts.elastic.co/packages/8.x/apt stable mai
n" | sudo tee /etc/apt/sources.list.d/elastic-8.x.list
deb [signed-by=/usr/share/keyrings/elasticsearch-keyring.gpg] https://artifacts.elastic.co/packages/8.x/apt stable main
root@ip-172-31-34-121:~# apt update
Hit:1 https://apt.grafana.com stable InRelease
Hit:2 https://download.docker.com/linux/ubuntu jammy InRelease
Hit:3 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy InRelease
Hit:4 https://artifacts.elastic.co/packages/7.x/apt stable InRelease
Ign:5 https://pkg.jenkins.io/debian-stable binary/ InRelease
Hit:6 https://pkg.jenkins.io/debian-stable binary/ Release
Hit:7 https://artifacts.elastic.co/packages/8.x/apt stable InRelease
Hit:8 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy-updates InRelease
Get:10 http://ports.ubuntu.com/ubuntu-ports jammy-security InRelease [129 kB]
Hit:11 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy-backports InRelease
Hit:9 https://prod-cdn.packages.k8s.io/repositories/istio/kubernetes:/core:/stable:/v1.30/deb InRelease
Hit:12 https://ppa.launchpadcontent.net/ansible/ansible/ubuntu jammy InRelease
Hit:13 https://ppa.launchpadcontent.net/pipewire-debian/pipewire-upstream/ubuntu jammy InRelease
Fetched 129 kB in 2s (69.7 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
152 packages can be upgraded. Run 'apt list --upgradable' to see them.
W: https://download.docker.com/linux/ubuntu/dists/jammy/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPREC
ATION section in apt-key(8) for details.
root@ip-172-31-34-121:~# vim /etc/logstash/conf.d/apache.conf
```

1.7 Copy and paste the following configurations

```
input {
  kafka {
    bootstrap_servers => "localhost:9092"
    topics => "apache"
  }
}
filter {
  grok {
    match => { "message" => "%{COMBINEDAPACHELOG}" }
  }
  date {
    match => [ "timestamp" , "dd/MMM/yyyy:HH:mm:ss Z" ]
  }
  geoip {
    source => "source_ip"
    target => "source_geo"
  }
}
output {
  elasticsearch {
    hosts => ["localhost:9200"]
    index => "logstash"
  }
}
```

```

root@ip-172-31-43-84:~# vim /etc/logstash/conf.d/apache.conf
root@ip-172-31-43-84:~# cat /etc/logstash/conf.d/apache.conf
input {
  kafka {
    bootstrap_servers => "localhost:9092"
    topics => "apache"
  }
}
filter {
  grok {
    match => { "message" => "%{COMBINEDAPACHELOG}" }
  }
  date {
    match => [ "timestamp" , "dd/MMM/yyyy:HH:mm:ss Z" ]
  }
  geoip {
    source => "source_ip"
    target => "source_geo"
  }
}
output {
  elasticsearch {
    hosts => ["localhost:9200"]
    index => "logstash"
  }
}
root@ip-172-31-43-84:~# █

```

1.8 Execute the following commands to install the Apache package:

apt install -y apache2

service apache2 start

service apache2 status

```

root@ip-172-31-43-84:~# apt install -y apache2
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  apache2-bin apache2-data apache2-utils libapr1 libaprutil1 libaprutil1-dbd-sqlite3 libaprutil1-ldap liblua5.3-0
Suggested packages:
  apache2-doc apache2-suexec-pristine | apache2-suexec-custom
The following NEW packages will be installed:
  apache2 apache2-bin apache2-data apache2-utils libapr1 libaprutil1 libaprutil1-dbd-sqlite3 libaprutil1-ldap liblua5.3-0
0 upgraded, 9 newly installed, 0 to remove and 267 not upgraded.
Need to get 2062 kB of archives.
After this operation, 8234 kB of additional disk space will be used.

```

```

root@ip-172-31-43-84:~# service apache2 start
root@ip-172-31-43-84:~# service apache2 status
● apache2.service - The Apache HTTP Server
   Loaded: loaded (/lib/systemd/system/apache2.service; enabled; vendor preset: enabled)
   Active: active (running) since Thu 2024-09-12 16:50:14 UTC; 29s ago
     Docs: https://httpd.apache.org/docs/2.4/
  Main PID: 15081 (apache2)
    Tasks: 55 (limit: 18808)
   Memory: 5.3M
      CPU: 35ms
   CGroup: /system.slice/apache2.service
           └─15081 /usr/sbin/apache2 -k start
             └─15083 /usr/sbin/apache2 -k start
               └─15084 /usr/sbin/apache2 -k start

Sep 12 16:50:14 ip-172-31-43-84 systemd[1]: Starting The Apache HTTP Server...
Sep 12 16:50:14 ip-172-31-43-84 systemd[1]: Started The Apache HTTP Server.
root@ip-172-31-43-84:~#

```

1.9 Execute the following command to install the Filebeat package:

apt -y install filebeat

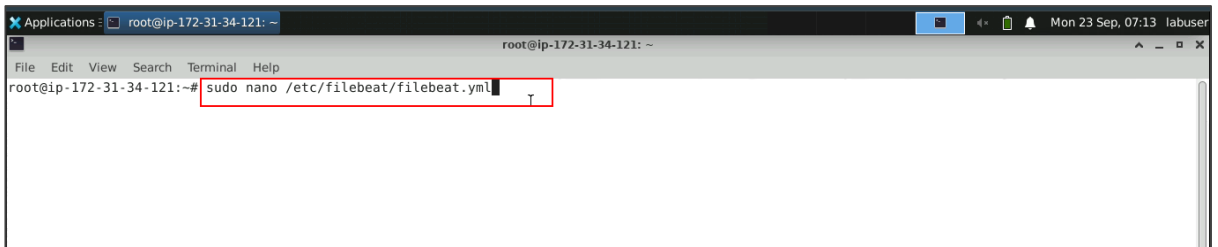
```

root@ip-172-31-43-84:~# apt -y install filebeat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following NEW packages will be installed:
  filebeat
0 upgraded, 1 newly installed, 0 to remove and 267 not upgraded.
Need to get 54.6 MB of archives.
After this operation, 201 MB of additional disk space will be used.
Get:1 https://artifacts.elastic.co/packages/8.x/apt stable/main amd64 filebeat amd64 8.15.1 [54.6 MB]
Fetched 54.6 MB in 1s (58.0 MB/s)
Selecting previously unselected package filebeat.
(Reading database ... 332039 files and directories currently installed.)
Preparing to unpack .../filebeat_8.15.1_amd64.deb ...
Unpacking filebeat (8.15.1) ...
Setting up filebeat (8.15.1) ...
Scanning processes...
Scanning linux images...

```

1.10 Run the command below to modify Filebeat so it can send logs to Logstash:

sudo nano /etc/filebeat/filebeat.yml



```

Applications: root@ip-172-31-34-121: ~
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# sudo nano /etc/filebeat/filebeat.yml

```

1.11 Update the **YAML** file using the following script:

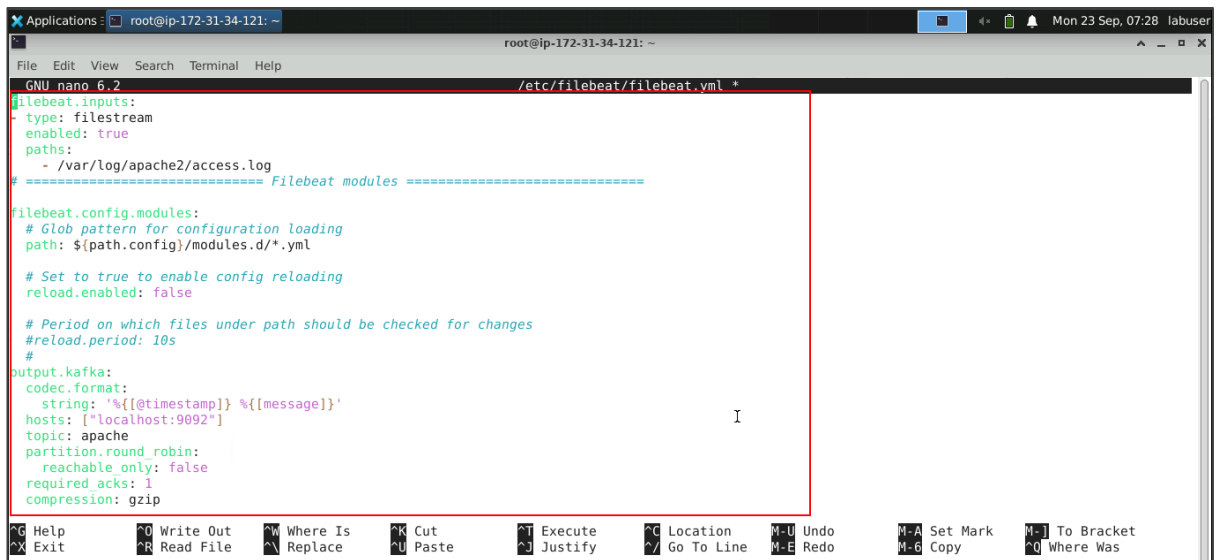
```
filebeat.inputs:
  - type: filestream
    enabled: true
    paths:
      - /var/log/apache2/access.log
# ===== Filebeat modules
=====

filebeat.config.modules:
  # Glob pattern for configuration loading
  path: ${path.config}/modules.d/*.yml

  # Set to true to enable config reloading
  reload.enabled: false

  # Period on which files under path should be checked for changes
  #reload.period: 10s
  #

output.kafka:
  codec.format:
    string: '%{[@timestamp]} %{[message]}'
  hosts: ["localhost:9092"]
  topic: apache
  partition.round_robin:
    reachable_only: false
  required_acks: 1
  compression: gzip
  max_message_bytes: 1000000
```

```
GNU nano 6.2 /etc/filebeat/filebeat.yml
filebeat.inputs:
- type: filestream
  enabled: true
  paths:
    - /var/log/apache2/access.log
# ===== Filebeat modules =====

filebeat.config.modules:
# Glob pattern for configuration loading
path: ${path.config}/modules.d/*.yml


# Set to true to enable config reloading
reload.enabled: false

# Period on which files under path should be checked for changes
#reload.period: 10s
#
output.kafka:
  codec.format:
    string: '%{[@timestamp]} %{[message]}'
  hosts: ["localhost:9092"]
  topic: apache
  partition.round_robin:
    reachable_only: false
  required_acks: 1
  compression: gzip
```

Step 2: Set up the Apache Kafka message broker

- 2.1 Execute the following command to install and set up the Apache Kafka message broker:

sudo apt-get install zookeeperd



```
root@ip-172-31-34-121:~# sudo nano /etc/filebeat/filebeat.yml
root@ip-172-31-34-121:~# sudo apt-get install zookeeperd
```

Note: Kafka uses ZooKeeper to maintain configuration information and synchronization.

```
Applications: root@ip-172-31-34-121: ~
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# sudo nano /etc/filebeat/filebeat.yml
root@ip-172-31-34-121:~# sudo apt-get install zookeeper
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  libjline-java liblog4j1.2-java libxerces2-java libxml-commons-external-java libxml-commons-resolver1.1-java libzookeeper-java zookeeper
Suggested packages:
  libjline-java-doc liblog4j1.2-java-doc libmail-java libxerces2-java-doc libxml-commons-resolver1.1-java-doc libzookeeper-java-doc
The following NEW packages will be installed:
  libjline-java liblog4j1.2-java libxerces2-java libxml-commons-external-java libxml-commons-resolver1.1-java libzookeeper-java zookeeper
  zookeeperd
0 upgraded, 8 newly installed, 0 to remove and 152 not upgraded.
Need to get 3778 kB of archives.
After this operation, 4548 kB of additional disk space will be used.
Do you want to continue? [Y/n] y
Get:1 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy/universe arm64 libjline-java all 1.0-3 [71.5 kB]
Get:2 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy/universe arm64 liblog4j1.2-java all 1.2.17-11 [439 kB]
Get:3 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy/universe arm64 libxml-commons-external-java all 1.4.01-5 [240 kB]
Get:4 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy/universe arm64 libxml-commons-resolver1.1-java all 1.2-11 [97.6 kB]
Get:5 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy/universe arm64 libxerces2-java all 2.12.1-1 [1437 kB]
Get:6 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy-updates/universe arm64 libzookeeper-java all 3.4.13-6ubuntu4.1 [1372 kB]
Get:7 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy-updates/universe arm64 zookeeper all 3.4.13-6ubuntu4.1 [112 kB]
Get:8 http://us-east-1.ec2.ports.ubuntu.com/ubuntu-ports jammy-updates/universe arm64 zookeeperd all 3.4.13-6ubuntu4.1 [8290 B]
Fetched 3778 kB in 3s (1408 kB/s)
```

- 2.2 Run the following commands to download and extract Kafka:
- ```
wget https://downloads.apache.org/kafka/3.8.0/kafka_2.12-3.8.0.tgz
tar -xvzf kafka_2.12-3.8.0.tgz
sudo cp -r kafka_2.12-3.8.0 /opt/kafka
```

```
Applications: root@ip-172-31-34-121: ~
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# wget https://downloads.apache.org/kafka/3.8.0/kafka_2.12-3.8.0.tgz
```

```
Applications: root@ip-172-31-34-121: ~
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# wget https://downloads.apache.org/kafka/3.8.0/kafka_2.12-3.8.0.tgz
--2024-09-23 07:32:47-- https://downloads.apache.org/kafka/3.8.0/kafka_2.12-3.8.0.tgz
Resolving downloads.apache.org (downloads.apache.org)... 88.99.208.237, 135.181.214.104, 2a01:4f9:3a:2c57::2, ...
Connecting to downloads.apache.org (downloads.apache.org)|88.99.208.237|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 120900366 (115M) [application/x-gzip]
Saving to: 'kafka_2.12-3.8.0.tgz'

kafka_2.12-3.8.0.tgz 100%[=====>] 115.30M 28.0MB/s in 5.3s

2024-09-23 07:32:53 (21.8 MB/s) - 'kafka_2.12-3.8.0.tgz' saved [120900366/120900366]

root@ip-172-31-34-121:~#
```

```
Applications: root@ip-172-31-34-121: ~
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# wget https://downloads.apache.org/kafka/3.8.0/kafka_2.12-3.8.0.tgz
--2024-09-23 07:32:47-- https://downloads.apache.org/kafka/3.8.0/kafka_2.12-3.8.0.tgz
Resolving downloads.apache.org (downloads.apache.org)... 88.99.208.237, 135.181.214.104, 2a01:4f9:3a:2c57::2, ...
Connecting to downloads.apache.org (downloads.apache.org)|88.99.208.237|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 120900366 (115M) [application/x-gzip]
Saving to: 'kafka_2.12-3.8.0.tgz'

kafka_2.12-3.8.0.tgz 100%[=====] 115.30M 28.0MB/s in 5.3s

2024-09-23 07:32:53 (21.8 MB/s) - 'kafka_2.12-3.8.0.tgz' saved [120900366/120900366]

root@ip-172-31-34-121:~# tar -xvzf kafka_2.12-3.8.0.tgz
```

```
Applications: root@ip-172-31-34-121: ~
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# sudo cp -r kafka_2.12-3.8.0 /opt/kafka
```

2.3 Run the following command to start Kafka:  
**sudo /opt/kafka/bin/kafka-server-start.sh /opt/kafka/config/server.properties**

```
Applications: root@ip-172-31-34-121: ~
root@ip-172-31-34-121: ~
File Edit View Search Terminal Help
root@ip-172-31-34-121:~# sudo cp -r kafka_2.12-3.8.0 /opt/kafka
root@ip-172-31-34-121:~# sudo /opt/kafka/bin/kafka-server-start.sh /opt/kafka/config/server.properties
```

You will see the following interface:

```
e may not be available. (org.apache.kafka.clients.NetworkClient)
[2024-09-02 15:31:08,034] INFO [Controller id=0, targetBrokerId=0] Client requested connection close from node 0 (org.apache.kafka.clients.NetworkClient)
[2024-09-02 15:31:08,078] INFO [/config/changes-event-process-thread]: Starting (kafka.common.ZkNodeChangeNotificationListener$ChangeEventProcessThread)
[2024-09-02 15:31:08,115] INFO [SocketServer listenerType=ZK_BROKER, nodeId=0] Enabling request processing. (kafka.network.SocketServer)
[2024-09-02 15:31:08,122] INFO [Awaiting socket connections on 0.0.0.0:9092. (kafka.network.DataPlaneAcceptor)
[2024-09-02 15:31:08,158] INFO [KafkaServer id=0] Start processing authorizer futures (kafka.server.KafkaServer)
[2024-09-02 15:31:08,159] INFO [KafkaServer id=0] End processing authorizer futures (kafka.server.KafkaServer)
[2024-09-02 15:31:08,159] INFO [KafkaServer id=0] Start processing enable request processing future (kafka.server.KafkaServer)
[2024-09-02 15:31:08,160] INFO [KafkaServer id=0] End processing enable request processing future (kafka.server.KafkaServer)
[2024-09-02 15:31:08,167] INFO Kafka version: 3.8.0 (org.apache.kafka.common.utils.AppInfoParser)
[2024-09-02 15:31:08,167] INFO Kafka commitId: 771b9576b08ecf5b (org.apache.kafka.common.utils.AppInfoParser)
[2024-09-02 15:31:08,167] INFO Kafka startTimeMs: 1725291068160 (org.apache.kafka.common.utils.AppInfoParser)
[2024-09-02 15:31:08,171] INFO [KafkaServer id=0] started (kafka.server.KafkaServer)
[2024-09-02 15:31:08,283] INFO [zk-broker-0-to-controller-forwarding-channel-manager]: Recorded new ZK controller, from now on will use node ip-172-31-36-251.ec2.intern
al:9092 (id: 0 rack: null) (kafka.server.NodeToControllerRequestThread)
[2024-09-02 15:31:08,339] INFO [zk-broker-0-to-controller-alter-partition-channel-manager]: Recorded new ZK controller, from now on will use node ip-172-31-36-251.ec2.i
nternal:9092 (id: 0 rack: null) (kafka.server.NodeToControllerRequestThread)
```

### Step 3: Create a topic for the Apache logs

- 3.1 Execute the following command in a new terminal tab to create a topic for the Apache logs:

```
/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --
replication-factor 1 --partitions 1 --topic apache
```

```
root@ip-172-31-36-251:~# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic apache
Created topic apache.
root@ip-172-31-36-251:~# █
```

The **Created topic apache** message is displayed as shown below:

```
root@ip-172-31-36-251:~# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic apache
Created topic apache.
root@ip-172-31-36-251:~# █
```

- 3.2 Run the commands below to start the Filebeat service, enable Filebeat at system startup, and check the status:

```
sudo systemctl start filebeat
sudo systemctl enable filebeat
sudo systemctl status filebeat
```

```
root@ip-172-31-43-84:~# sudo systemctl start filebeat
root@ip-172-31-43-84:~#
root@ip-172-31-43-84:~# sudo systemctl enable filebeat
Synchronizing state of filebeat.service with SysV service script with /lib/systemd/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable filebeat
Created symlink /etc/systemd/system/multi-user.target.wants/filebeat.service → /lib/systemd/system/filebeat.service.
root@ip-172-31-43-84:~#
root@ip-172-31-43-84:~# sudo systemctl status filebeat
● filebeat.service - Filebeat sends log files to Logstash or directly to Elasticsearch.
 Loaded: loaded (/lib/systemd/system/filebeat.service; enabled; vendor preset: enabled)
 Active: active (running) since Thu 2024-09-12 16:44:39 UTC; 14s ago
 Docs: https://www.elastic.co/beats/filebeat
 Main PID: 13336 (filebeat)
 Tasks: 9 (limit: 18808)
 Memory: 36.8M
 CPU: 110ms
 CGroup: /system.slice/filebeat.service
 └─13336 /usr/share/filebeat/bin/filebeat --environment systemd -c /etc/filebeat/filebeat.yml --path.home /usr/share/filebeat

Sep 12 16:44:39 ip-172-31-43-84 systemd[1]: Started Filebeat sends log files to Logstash or directly to Elasticsearch..
root@ip-172-31-43-84:~# █
```

3.3 Run the following command to verify that Elasticsearch is receiving the Filebeat data log:

```
curl -X GET "localhost:9200/_cat/indices?v" | grep logstash
```

```
root@ip-172-31-21-55:/var/log/filebeat# curl -X GET "localhost:9200/_cat/indices?v" | grep logstash
% Total % Received % Xferd Average Speed Time Time Time Current
 Dload Upload Total Spent Left Speed
100 2975 0 2975 0 0 442k 0 --:--:-- --:--:-- --:--:-- 484k
yellow open logstash NakDRkMKTICd-xAaAA4gHA 1 1 3 0 20.4kb 20.4kb
20.4kb
root@ip-172-31-21-55:/var/log/filebeat#
```

The displayed output is as shown below:

```
root@ip-172-31-21-55:/var/log/filebeat# curl -X GET "localhost:9200/_cat/indices?v" | grep logstash
% Total % Received % Xferd Average Speed Time Time Time Current
 Dload Upload Total Spent Left Speed
100 2975 0 2975 0 0 442k 0 --:--:-- --:--:-- --:--:-- 484k
yellow open logstash NakDRkMKTICd-xAaAA4gHA 1 1 3 0 20.4kb 20.4kb
20.4kb
root@ip-172-31-21-55:/var/log/filebeat#
```

By following these steps, you have successfully integrated Kafka with ELK Stack to buffer and process log surges, protecting Logstash and Elasticsearch during high-volume events for improved reliability and scalable log management.