

Lesson-End Project

Implementing Jenkins Logging using ELK Stack

Project agenda: To create Jenkins logging integration with ELK Stack for centralized log collection and analysis

Description: You have been given a task to onboard the Jenkins application for implementing CI/CD pipelines. Once pipelines are created you need to ensure Jenkins is working fine as expected. So, you will integrate Jenkins logging on ELK stack so that developers don't need root access to monitor Jenkins logs.

Tools required: Jenkins, ELK Stack

Prerequisites: None

Expected deliverables: A fully integrated Jenkins logging system with the ELK Stack for centralized log collection and analysis. This system should enable efficient visualization and monitoring of Jenkins logs through Kibana, without requiring root access for developers to monitor the logs

Steps to be followed:

1. Configure file beat to collect Jenkins logs
2. Configure the Kibana visualization tool to check Jenkins logs

Step 1: Configure file beat to collect Jenkins logs

1.1 Run the following command on the terminal to download and install Filebeat:

sudo su

apt-get install filebeat

```
labuser@ip-172-31-15-221:~$ sudo su
root@ip-172-31-15-221:/home/labuser# apt-get install filebeat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following NEW packages will be installed:
  filebeat
0 upgraded, 1 newly installed, 0 to remove and 152 not upgraded.
Need to get 47.9 MB of archives.
After this operation, 192 MB of additional disk space will be used.
Get:1 https://artifacts.elastic.co/packages/8.x/apt stable/main arm64 filebeat arm64 8.15.1 [47.9 MB]
Fetched 47.9 MB in 7s (7199 kB/s)
Selecting previously unselected package filebeat.
(Reading database ... 293605 files and directories currently installed.)
Preparing to unpack .../filebeat_8.15.1_arm64.deb ...
Unpacking filebeat (8.15.1) ...
Setting up filebeat (8.15.1) ...
Scanning processes...
Scanning linux images...
```

1.2 Execute the below command to modify Filebeat, so that it can send logs to Logstash

nano /etc/filebeat/filebeat.yml

```
No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
root@ip-172-31-15-221:/home/labuser# nano /etc/filebeat/filebeat.yml
```

1.3 Add the below code into the file:

```
filebeat.inputs:
- type: filestream
  paths:
    - "/var/log/jenkins/jenkins*.log"
  enabled: true

# ===== Elasticsearch template setting =====

setup.template.settings:
  index.number_of_shards: 1
  #index.codec: best_compression
  #_source.enabled: false

output.elasticsearch:
  hosts: ["localhost:9200"]
```

```
filebeat.inputs:
- type: filestream
  paths:
    - "/var/log/jenkins/jenkins*.log"
  enabled: true

# ===== Elasticsearch template setting =====

setup.template.settings:
  index.number_of_shards: 1
  #index.codec: best_compression
  #_source.enabled: false

output.elasticsearch:
  hosts: ["localhost:9200"]
```

1.4 Run the below command to open Service configuration using the vi command and uncomment the below code line as shown in the screenshot:

vi /lib/systemd/system/jenkins.service

```
No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
root@ip-172-31-15-221:/home/labuser# nano /etc/filebeat/filebeat.yml
root@ip-172-31-15-221:/home/labuser# vi /lib/systemd/system/jenkins.service
```

```
# Directory where Jenkins stores its configuration and workspaces
Environment="JENKINS_HOME=/var/lib/jenkins"
WorkingDirectory=/var/lib/jenkins

# Location of the Jenkins WAR
#Environment="JENKINS_WAR=/usr/share/java/jenkins.war"

# Location of the exploded WAR
Environment="JENKINS_WEBROOT=%C/jenkins/war"

# Location of the Jenkins log. By default, systemd-journald(8) is used.
Environment="JENKINS_LOG=%L/jenkins/jenkins.log"

# The Java home directory. When left empty, JENKINS_JAVA_CMD and PATH are consulted.
#Environment="JAVA_HOME=/usr/lib/jvm/java-17-openjdk-amd64"
```

Note: press esc and enter **:wq** to save and exit the file.

1.5 Once the configuration is saved run the below command to restart the Jenkins:

systemctl daemon-reload

service jenkins restart

```
No VM guests are running outdated hypervisor (qemu) binaries on this host.
root@ip-172-31-15-221:/home/labuser# nano /etc/filebeat/filebeat.yml
root@ip-172-31-15-221:/home/labuser# vi /lib/systemd/system/jenkins.service
root@ip-172-31-15-221:/home/labuser# systemctl daemon-reload
root@ip-172-31-15-221:/home/labuser# service jenkins restart
```

1.6 Run the below commands to restart the configuration file:

sudo systemctl start filebeat

sudo systemctl enable filebeat

sudo systemctl status filebeat

```
root@ip-172-31-15-221:/home/labuser# sudo systemctl start filebeat
root@ip-172-31-15-221:/home/labuser# sudo systemctl enable filebeat
Synchronizing state of filebeat.service with SysV service script with /lib/systemd/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable filebeat
Created symlink /etc/systemd/system/multi-user.target.wants/filebeat.service → /lib/systemd/system/filebeat.service.
root@ip-172-31-15-221:/home/labuser# sudo systemctl status filebeat
● filebeat.service - Filebeat sends log files to Logstash or directly to Elasticsearch.
   Loaded: loaded (/lib/systemd/system/filebeat.service; enabled; vendor preset: enabled)
   Active: active (running) since Wed 2024-09-25 09:49:23 UTC; 27s ago
     Docs: https://www.elastic.co/beats/filebeat
    Main PID: 10105 (filebeat)
      Tasks: 8 (limit: 9361)
     Memory: 64.1M
        CPU: 726ms
    CGroup: /system.slice/filebeat.service
            └─10105 /usr/share/filebeat/bin/filebeat --environment systemd -c /etc/filebeat/filebeat.yml --path.home /usr/share/filebeat --path.conf

Sep 25 09:49:23 ip-172-31-15-221 systemd[1]: Started Filebeat sends log files to Logstash or directly to Elasticsearch..
```

1.7 Execute the below command to verify that Elasticsearch is receiving the filebeat data log:

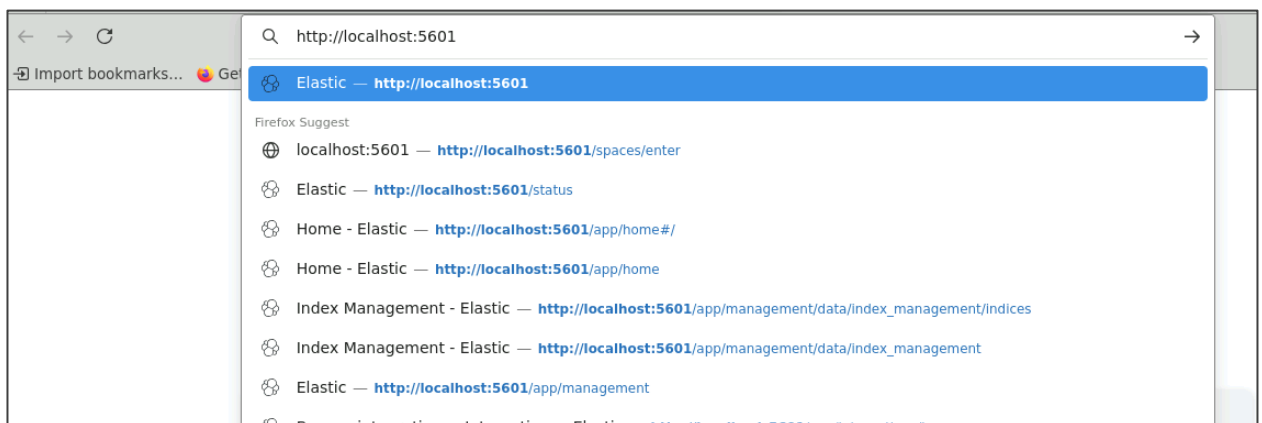
curl -XGET http://localhost:9200/_cat/indices?v | grep filebeat

```
root@ip-172-31-15-221:/home/labuser# curl -XGET http://localhost:9200/_cat/indices?v | grep filebeat
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           %             Dload  Upload  Total   Spent    Left   Speed
100 3500    0 3500    0     0  32134      0 --:--:-- --:--:-- --:--:-- 32407
yellow open .ds-filebeat-8.15.1-2024.09.25-000001      sMIh2u9mQa6IaAFsAj5xbg 1 1
26.8kb      26.8kb
root@ip-172-31-15-221:/home/labuser#
```

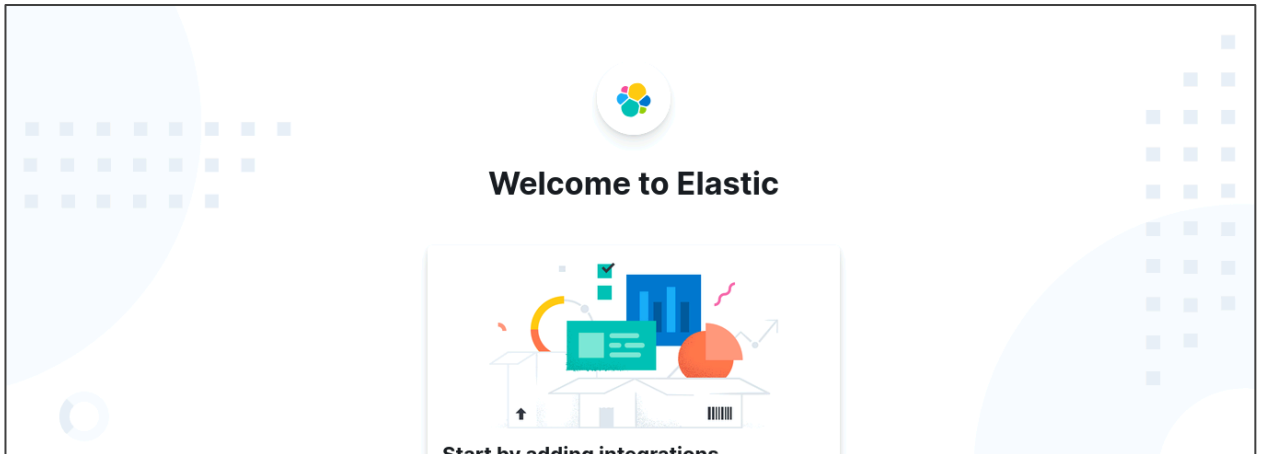
Step 2: Configure Kibana visualization tool to check Jenkins logs

2.1 Enter the below URL to access the Kibana web interface to visualized logs collected by filebeat

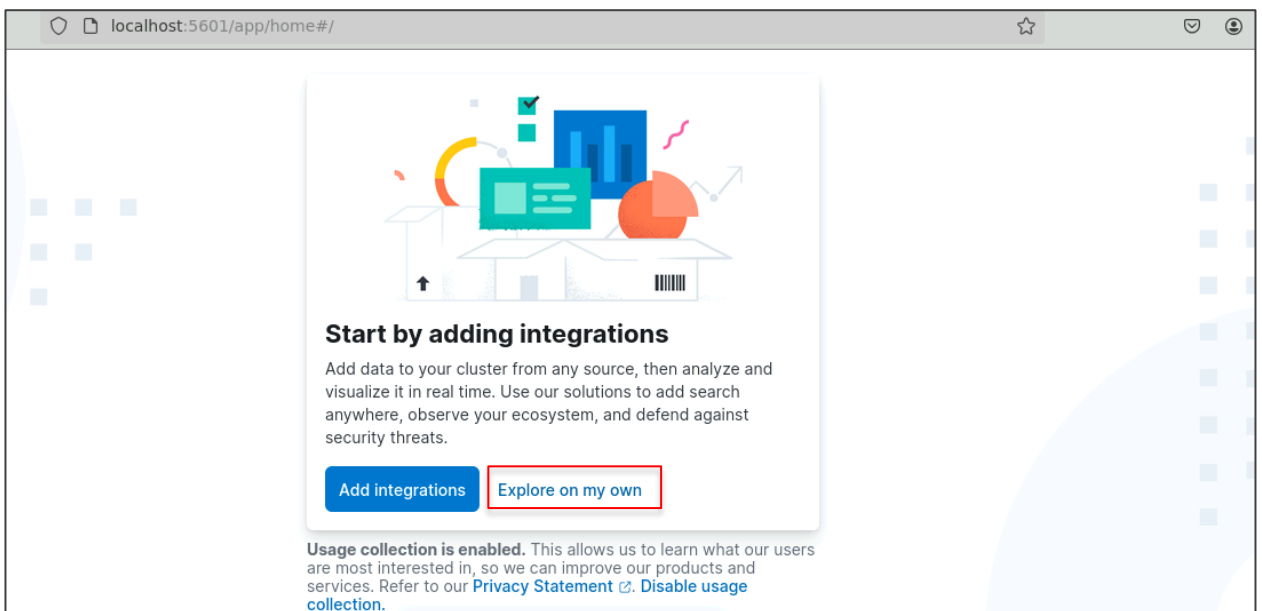
<http://localhost:5601>



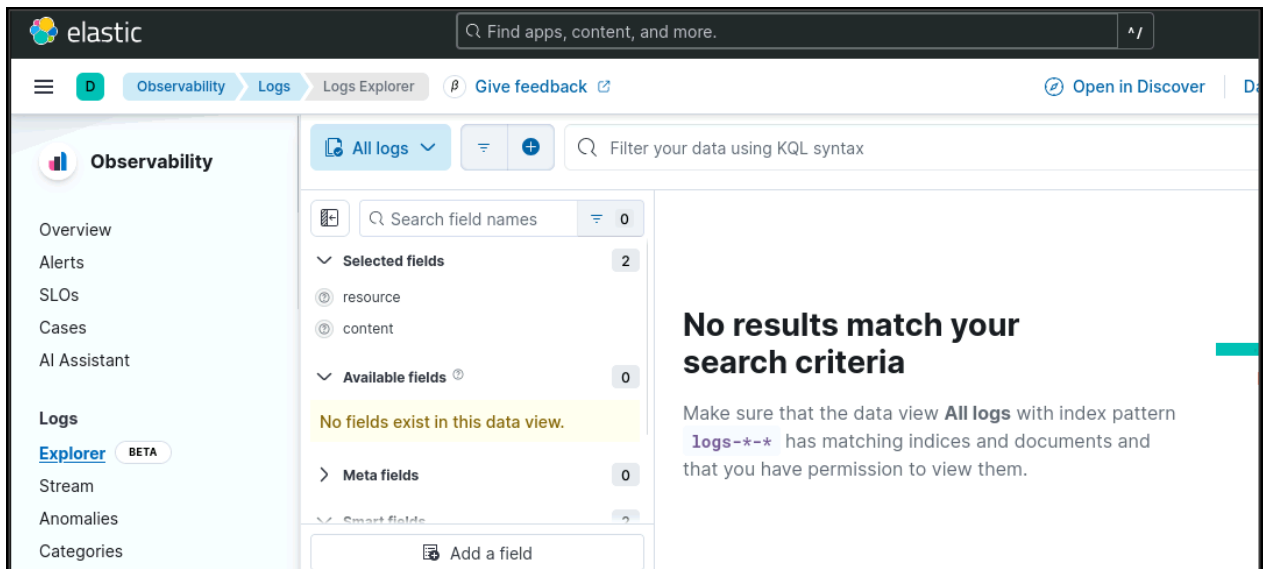
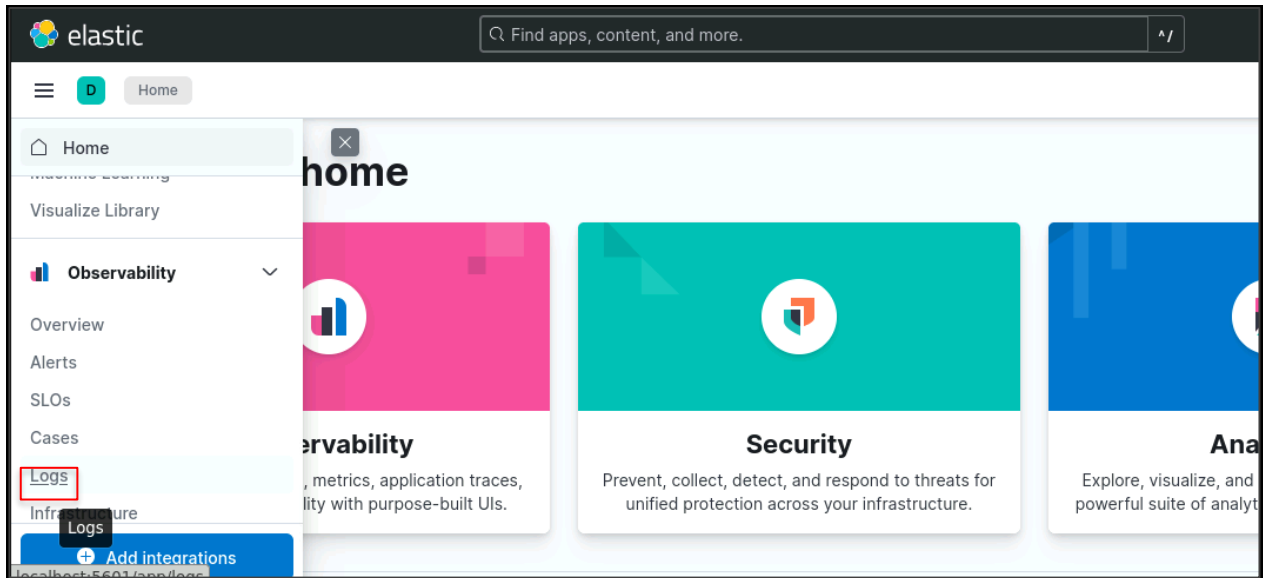
You can see the below screen:



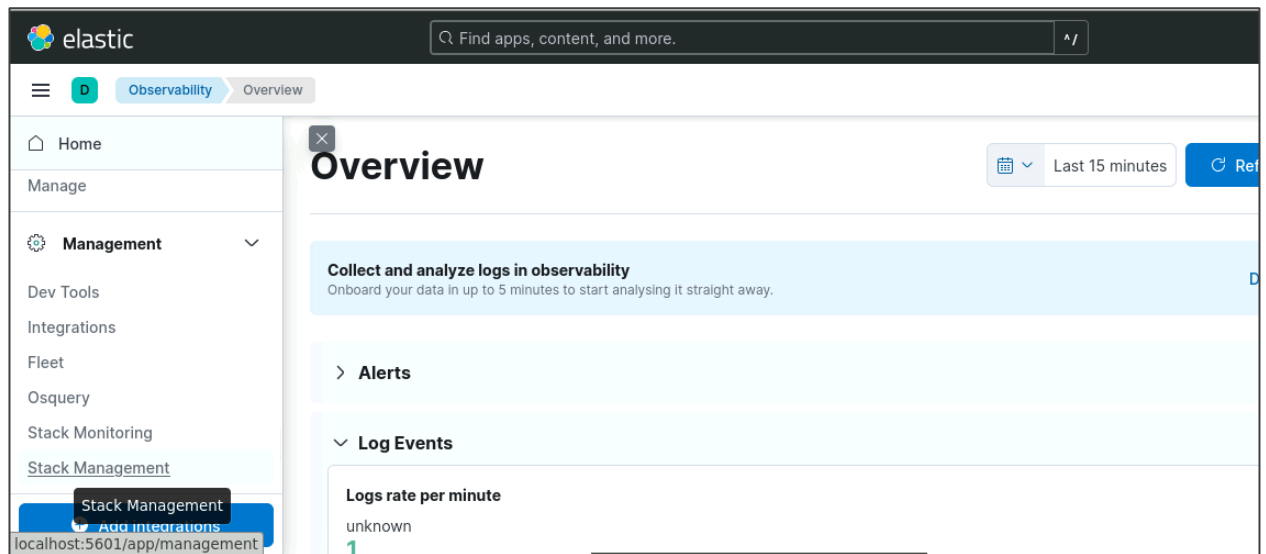
2.2 Click on **Explore on my own**



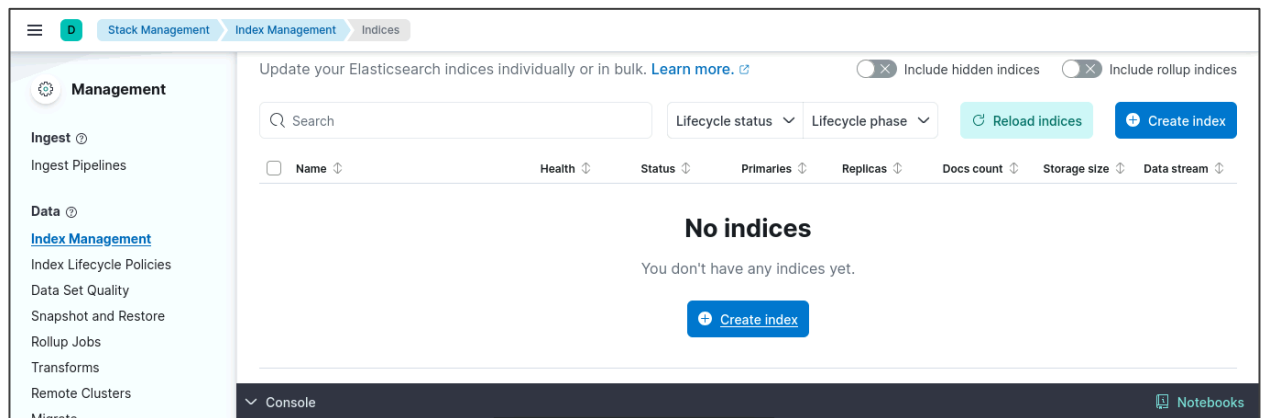
2.3 Click on Home, click **Logs** to check the system logs in Kibana



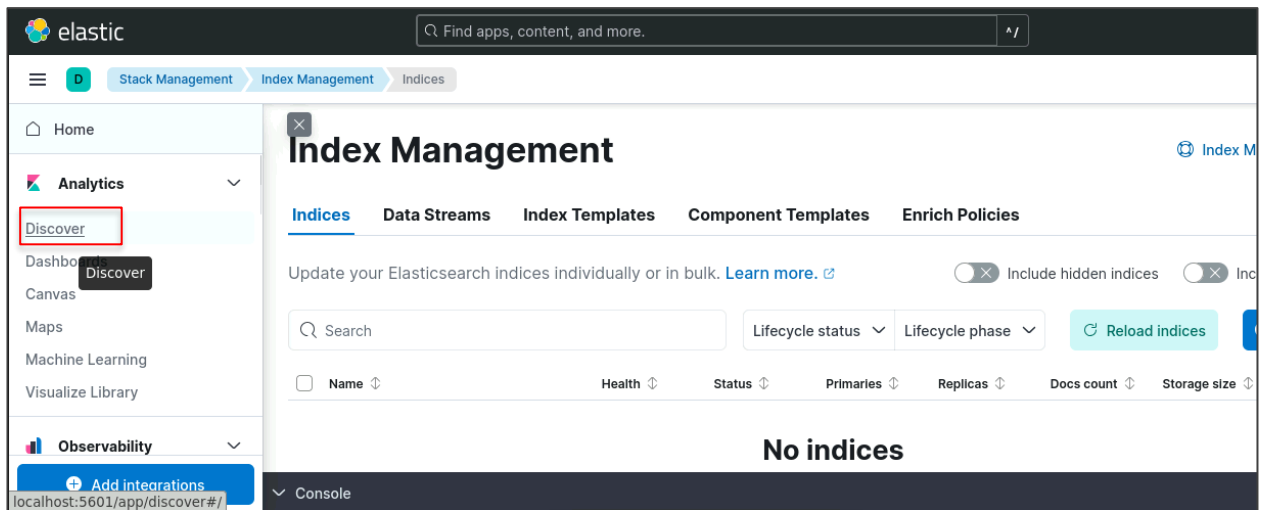
2.4 Now navigate back to **Home** and click on **Stack Management**



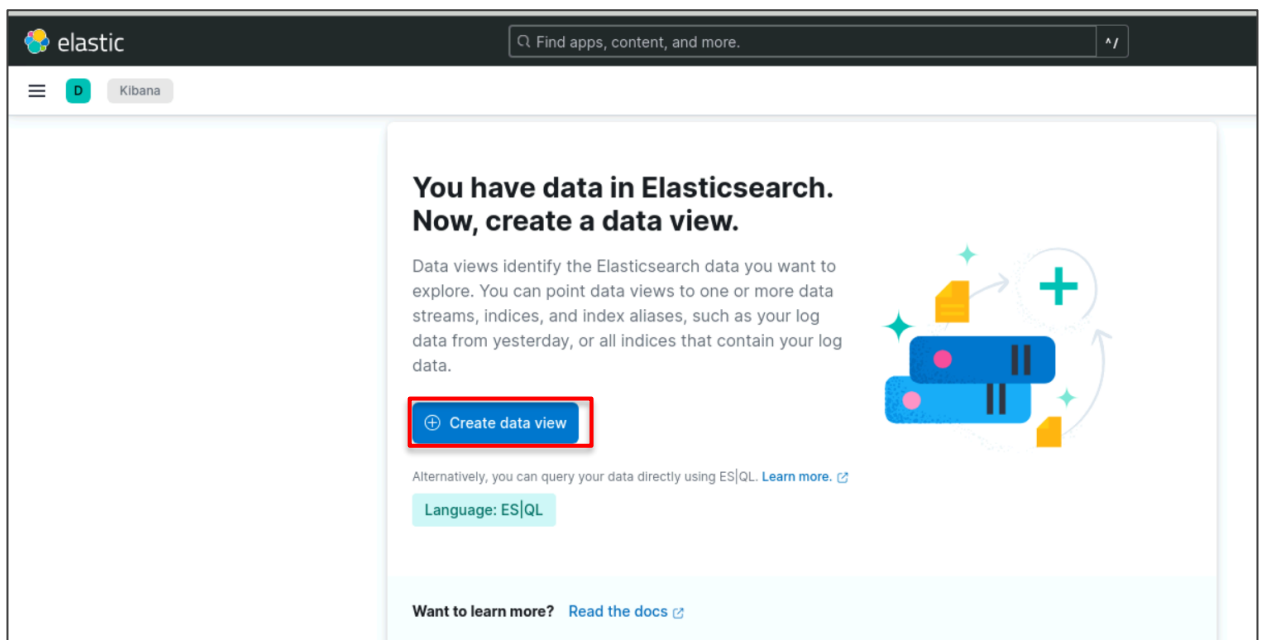
2.5 Click on **Index Management**



2.6 Under **Analytics** click on **Discover**



2.7 Click on **Create data view** then fill the details as shown in the screenshot and click on **Save data view to Kibana**



Create data view

Name

jenkins

Index pattern

filebeat*

Timestamp field

@timestamp

Select a timestamp field for use with the global time filter.

Show advanced settings

Close

Save data view to Kibana

Your index pattern matches 1 source.

All sources

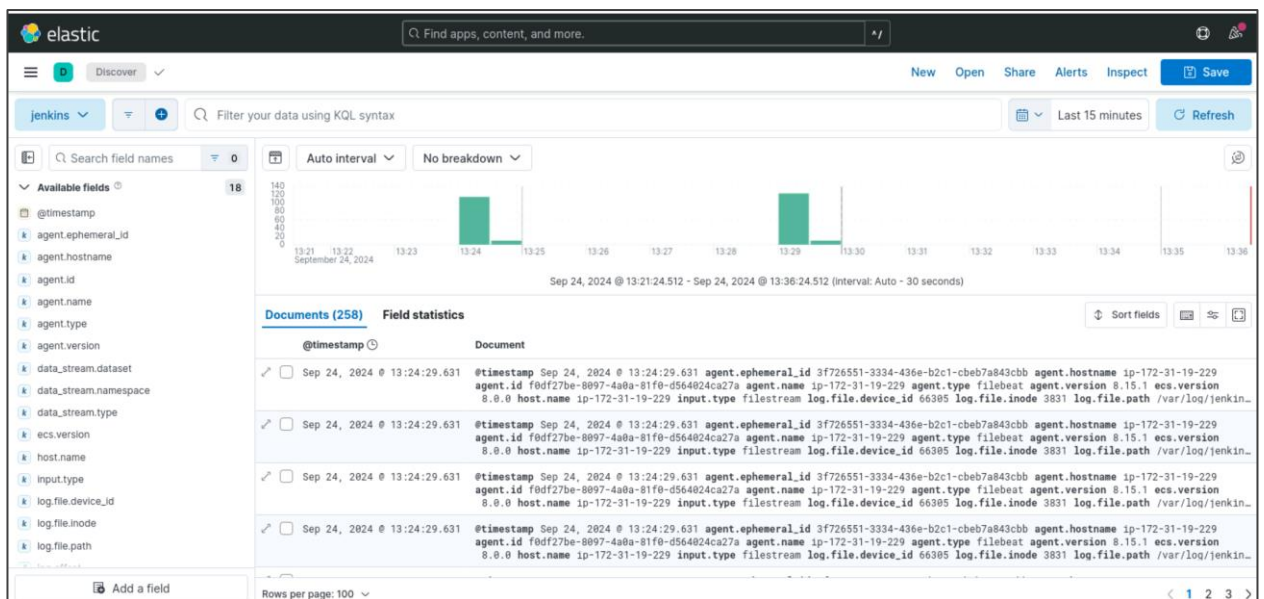
Matching sources

filebeat-8.15.1

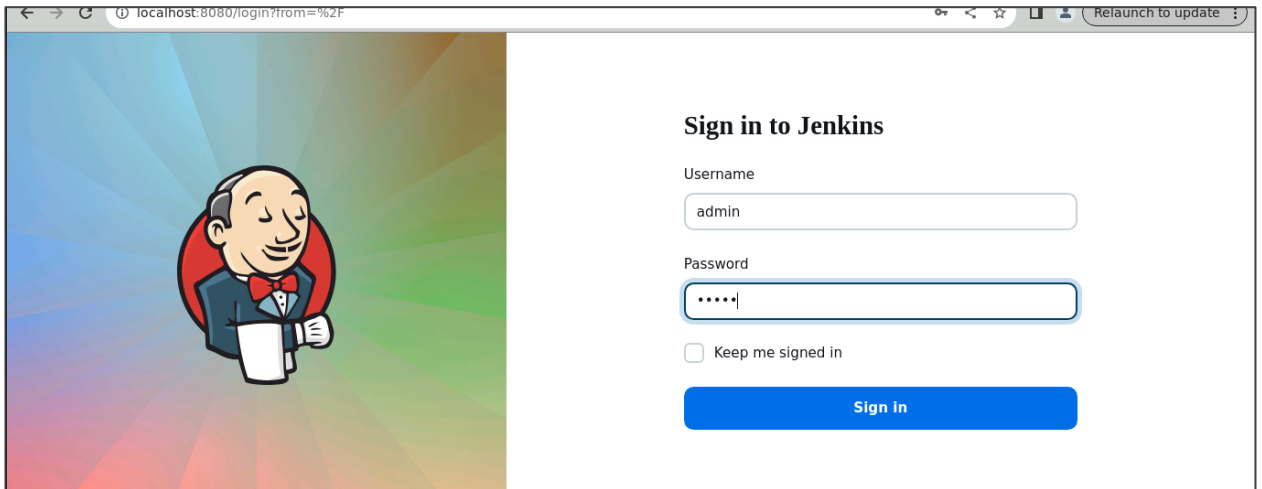
Data stream

Rows per page: 10

You can visualize the logs as shown below:



2.8 Navigate to Jenkins using the below **URL**:
localhost:8080



localhost:8080/login?from=%2F

Sign in to Jenkins

Username
admin

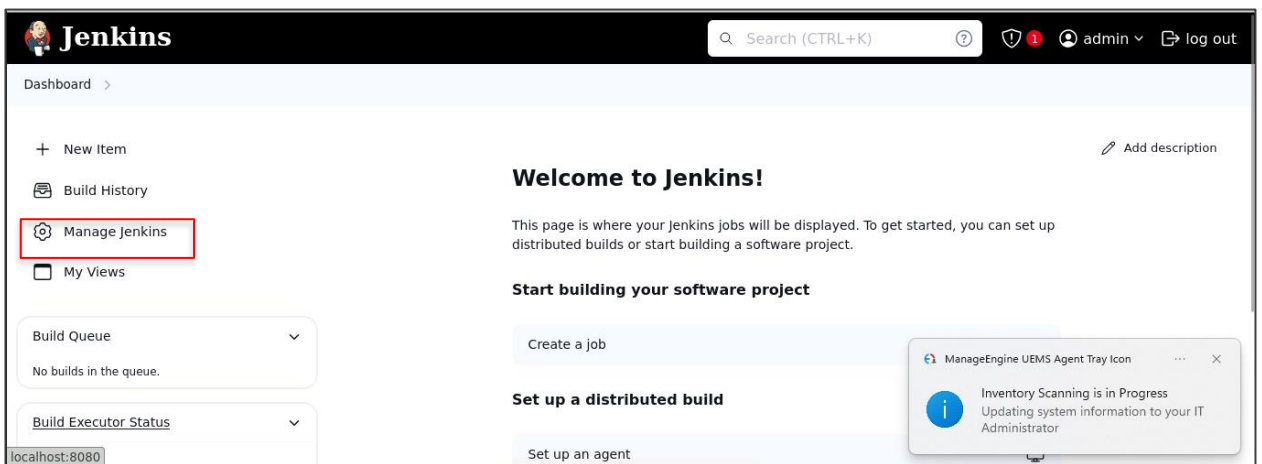
Password
.....

☐ Keep me signed in

Sign in

Note: Username: admin and Password: Root123\$

2.9 Click on **Manage Jenkins**



Jenkins

Search (CTRL+K)

admin log out

Dashboard >

- + New Item
- Build History
- Manage Jenkins**
- My Views

Build Queue
No builds in the queue.

Build Executor Status

localhost:8080

Welcome to Jenkins!

This page is where your Jenkins jobs will be displayed. To get started, you can set up distributed builds or start building a software project.

Start building your software project

Create a job

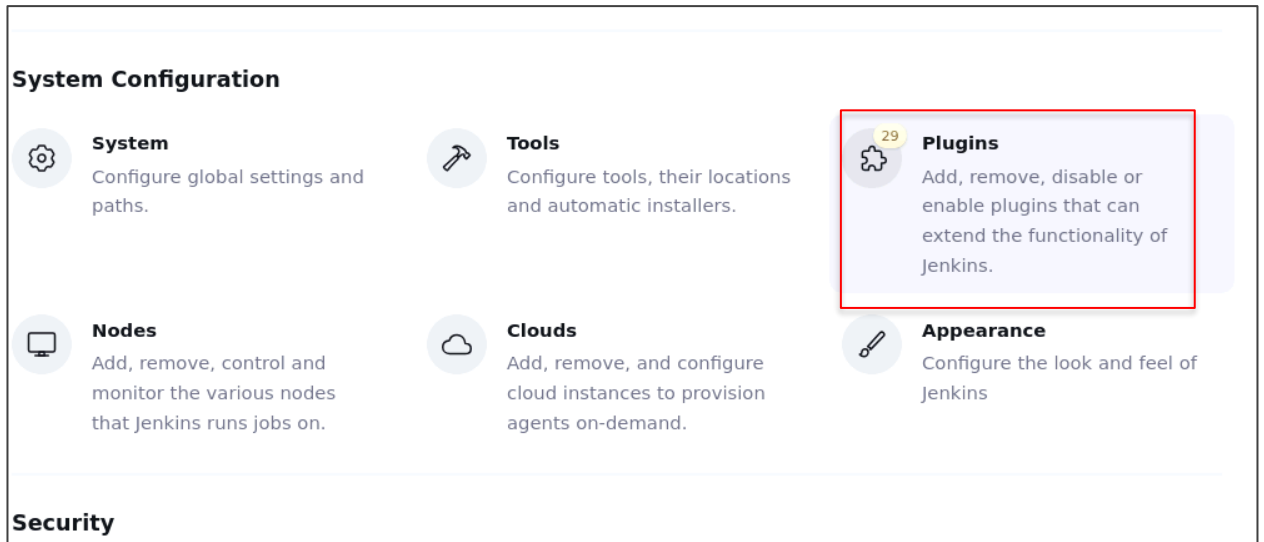
Set up a distributed build

Set up an agent

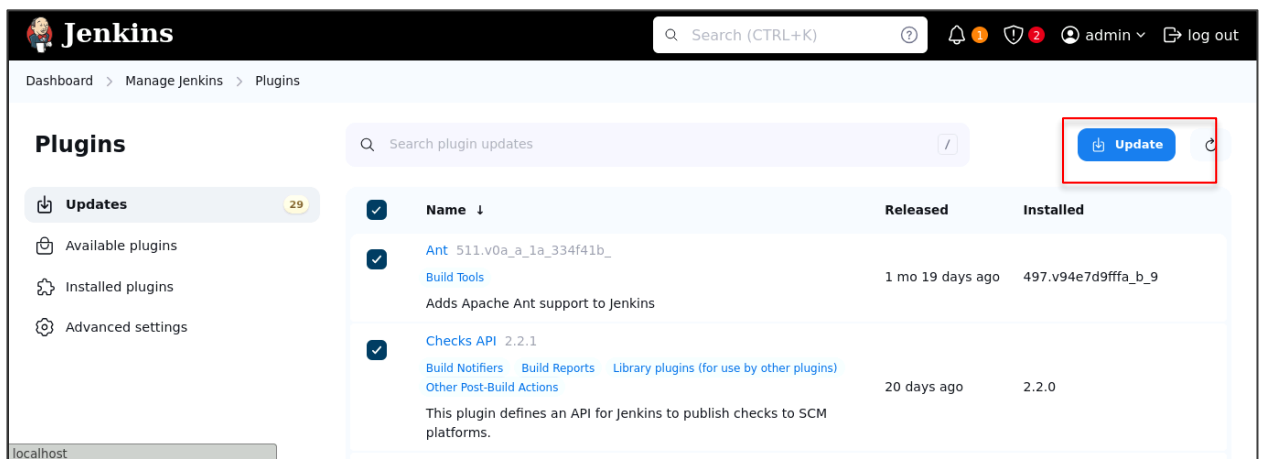
ManagementEngine UEMS Agent Tray Icon

Inventory Scanning is in Progress
Updating system information to your IT Administrator

2.10 Click on the **Plugins**

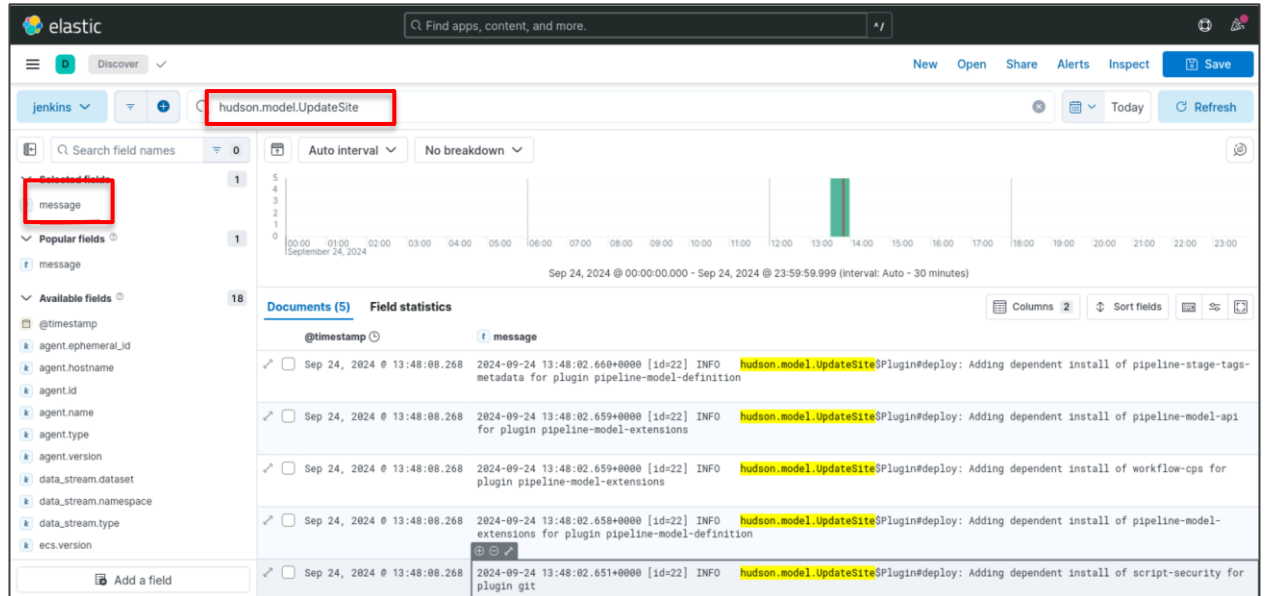


2.11 Select all the plugins and click on **Update**



- 2.12 Navigate back to Kibana and search for the below keyword to see recent log entries regarding plugin upgrade activity on Jenkins. Also, select the **message** field to get only log entries in Kibana without any metadata.

Keyword: hudson.model.UpdateSite



By following these steps, you have successfully created Jenkins logging integration with ELK Stack for centralized log collection and analysis.